



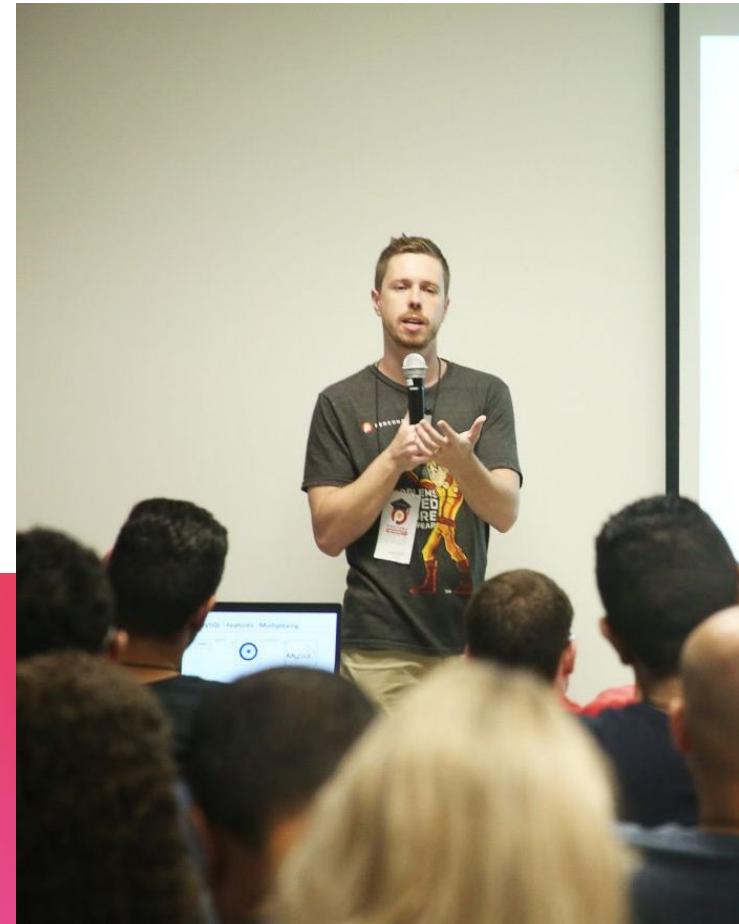
Troubleshooting MySQL from a MySQL Developer Perspective

Marcelo Altmann
Software Developer - Percona
Percona Live Austin / May 2022



Marcelo Altmann

- Software Developer @ Percona
 - Joined in 2016 as Senior Support Engineer
 - PXB / PS / PXC
 - Member Bugs Escalation Committee



Software bug

Software bug

A software bug is an error, flaw or fault in computer software that causes it to produce an incorrect or unexpected result, or to behave in unintended ways - [wikipedia](#)

- Wrong query results
- Performance degradations
- Security flaw
- Crash

Reproducible test case

The history of a dev



Reproducible test case

- Been able to reproduce a bug is key
- Projects have their own test framework
- Used to ensure we don't introduce regressions.
- MySQL MTR

- Start mysql instance:

```
./mtr --start alias &  
mysql --defaults-file=var/my.cnf
```

- Interactive GDB session:

```
./mtr --manual-gdb alias  
gdb -cd ./ -x ./var/tmp/gdbinit.mysql.1 ..../bin/mysql
```

Reproducible test case

- Recompile with [Debug synchronization](#) :

```
open_tables(...)  
DEBUG_SYNC(thd, "after_open_tables");  
lock_tables(...)  
  
--connection conn1  
SET DEBUG_SYNC= 'after_open_tables SIGNAL opened WAIT_FOR flushed';  
send INSERT INTO t1 VALUES(1);  
--connection conn2  
SET DEBUG_SYNC= 'now WAIT_FOR opened';  
SET DEBUG_SYNC= 'after_abort_locks SIGNAL flushed';  
FLUSH TABLE t1;
```

Reproducible test case

- Recompile with [Test Faults Macros](#) :

```
somefile.cc
```

```
DBUG_EXECUTE_IF("some_keyword", {
    some_function();
    some_variable++;
});
```

```
mysql> SET debug = '+d,some_keyword';
```

Stack traces

Crash fingerprint

Stack Traces - Signals

- A way of software / kernel / user to communicate with a process.
- Program will interpret and handle it - mysqld.cc my_init_signals()

```
/* ISO C99 signals. */
#define SIGINT          2      /* Interactive attention signal. */
#define SIGILL          4      /* Illegal instruction. */
#define SIGABRT         6      /* Abnormal termination. */
#define SIGFPE          8      /* Erroneous arithmetic operation. */
#define SIGSEGV         11     /* Invalid access to storage. */
#define SIGTERM         15     /* Termination request. */
```

Stack Traces - Signals

- SIGABRT / 6 - Code Assertions
 - ut_a(condition) / ut_ad(condition)
 - Critical points of code - continuing might cause damage
- SIGSEGV / 11 - Segmentation Fault
 - Memory access violation
 - Tried to access restricted memory area
- SIGTERM / 15 - Terminate
 - shutdown

Stack Traces

19:03:30 UTC - mysqld got signal 6 ;

Most likely, you have hit a bug, but this error can also be caused by malfunctioning hardware.

1

Build ID: ae15f57aa1c932e10b043d5f3e353ef09fe0e724

Server Version: 8.0.27-18-debug Percona Server (GPL), Revision 63a5835a303-Debug

2

Thread pointer: 0x7fe1f8001040

Attempting backtrace. You can use the following information to find out where mysqld died. If you see no messages after this, something went terribly wrong...

stack_bottom = 7fe2b02ebbd8 thread_stack 0x100000

mysqld(my_print_stacktrace(unsigned char const*, unsigned long)+0x55) [0x5579240abc6c]

mysqld(handle_fatal_signal+0x2e0) [0x557922e6e3ec]

...

mysqld(ut_dbg_assertion_failed(char const*, char const*, unsigned long)+0x1e3) [0x55792457bce8]

mysqld(lock_rec_convert_active_impl_to_expl(. . .) [0x5579242ffa15]

mysqld(row_convert_impl_to_expl_if_needed(btr_cur_t*, undo_node_t*)+0x143) [0x55792449ede2]

mysqld(+0x5586364) [0x5579248b0364]

mysqld(+0x55866ba) [0x5579248b06ba]

mysqld(+0x5588323) [0x5579248b2323]

3

Trying to get some variables.

Some pointers may be invalid and cause the dump to abort.

Query (7fe1f800a888): UPDATE t1 SET u=oo,v=ef WHERE id=10

Connection ID (thread ID): 11

Status: NOT_KILLED

4

Divide and conquer

Page Cleaner bug



Divide and conquer

- Goal is to establish when this started to happen
- Remove as much as unknown as you can
 - Issue is happening on latest version?
 - Issue is happening on latest version minus 1 / 2 / 3 / ... versions?
 - Issue is happening on previous major version (8.0 -> 5.7) ?
- (PXB / PXC) - Is the issue specific to the product?

Divide and conquer

- [PXB-2742](#) as example:
-



Percona XtraBackup / [PXB-2742](#)

PXB crashes during incremental backup prepare of partitioned tables: Assertion failure:
`buf0flu.cc:3567:UT_LIST_GET_LEN(buf_pool->flush_list) == 0`

- Start Point - partitions -> full backup -> inc backup (one or many) -> prepare full -> prepare inc (one or many) -> crash (sometimes)
- partitions ? = full backup -> inc backup (one or many) -> prepare full -> prepare inc (one or many) -> crash (sometimes)

Divide and conquer

- Incrementals ? = full backup -> prepare full -> crash (sometimes)
- Investigate (prepare full):
 - a. Innodb Master Thread Doing idle task
 - b. Merging Insert Buffer at full io capacity
 - Ask the pages to be read in to BP (async IO) (IO_BUF_READ)
 - c. IO Read Thread read the page
 - Merge Ibuf changes - Add page to Flush list
 - d. Shutdown
 - e. IO Read Thread complete the read/ibuf merge of page (IO_BUF_NONE)

Divide and conquer

- Can I reproduce the same on server?
- DBUG_EXECUTE_IF - shutdown - full ibuf merge
- Yes !!! 8.0 & 5.7 PS and upstream affected.
- [PS-8174](#) / [#107069](#) Crash -> private =(
- From a complex set of multiple variables (partition, xtrabackup, multiple incremental) to "simple" server bug.

Regression = git bisect

Wrong query result bug



PERCONA LIVE



Regression = git bisect

- [PS-7019 / #99398](#)
- Works ok on 8.0.19
- Does NOT work on 8.0.20

Regression = git bisect

```
mysql> SELECT * FROM t1;
+-----+-----+
| t1_id | t2_id |
+-----+-----+
|     1 |  1000 |
|     2 |      5 |
+-----+-----+
2 rows in set (0,00 sec)
```

```
mysql> SELECT * FROM t2;
+-----+-----+
| t2_id | is_active |
+-----+-----+
|     2 |         1 |
|     3 |         0 |
| 1000 |         1 |
+-----+-----+
3 rows in set (0,00 sec)
```

```
8.0.19> SELECT t1.*, t2.t2_id FROM t1 LEFT JOIN
t2 ON (t1.t2_id = t2.t2_id) GROUP BY t1_id;
+-----+-----+-----+
| t1_id | t2_id | t2_id |
+-----+-----+-----+
|     1 |  1000 |  1000 |
|     2 |      5 |    NULL |
+-----+-----+
2 rows in set (0,00 sec)
```

```
8.0.20> SELECT t1.*, t2.t2_id FROM t1 LEFT JOIN
t2 ON (t1.t2_id = t2.t2_id) GROUP BY t1_id;
+-----+-----+-----+
| t1_id | t2_id | t2_id |
+-----+-----+-----+
|     1 |  1000 |    NULL |
|     2 |      5 |    NULL |
+-----+-----+
2 rows in set (0,00 sec)
```

Regression = git bisect

- [Finding a Regression in MySQL Source Code: A Case Study](#)

- Lines (737+K):

```
git diff mysql-8.0.19..mysql-8.0.20 | wc -l  
737454
```

- Files (~4.5K):

```
git diff mysql-8.0.19..mysql-8.0.20 --name-only | wc -l  
4495
```

- Commits (~2K):

```
git log mysql-8.0.19..mysql-8.0.20 --pretty=oneline | wc -l  
1966
```

Regression = git bisect

Start Bad ← Commit 20
.....
Commit 15
.....
Current commit to test ← Commit 10
.....
Commit 5
.....
Start Good ← Commit 1

Commit 10
Commit 9
.....
Commit 7
.....
Commit 5
.....
Commit 3
.....
Commit 1

Commit 10 still has the problem (bad)

Regression = git bisect

Commit 10

Commit 9

...

Commit 7

Commit 5 is
good

Commit 5

...

Commit 3

...

Commit 1

Commit 10

Commit 9

Commit 8

Commit 7

Commit 6

Commit 5



PERCONA LIVE

Regression = git bisect

Commit 10

Commit 9

Commit 8 Commit 7 is Commit 7
bad

Commit 7



Commit 6

Commit 6

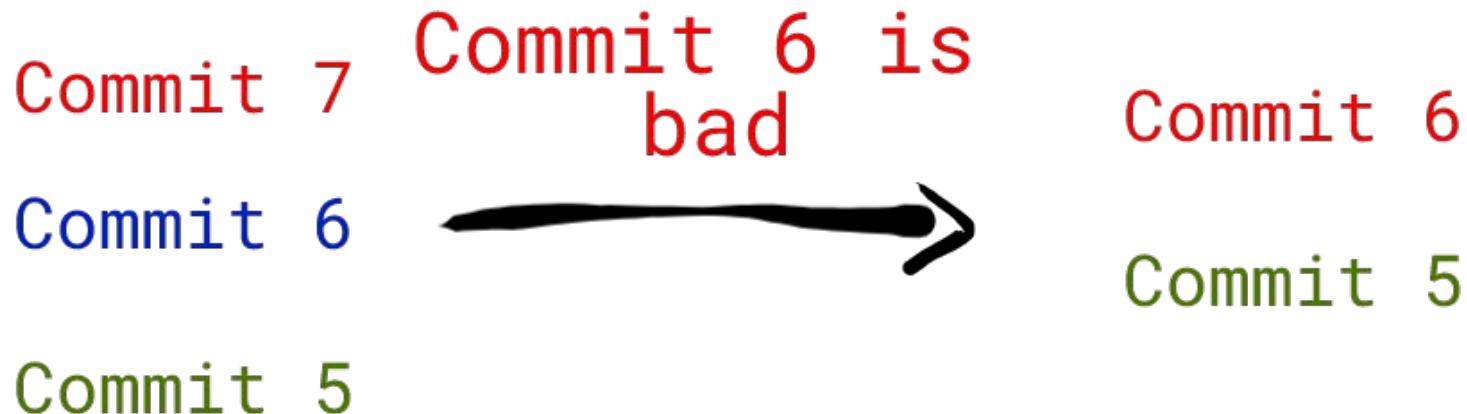
Commit 5

Commit 5



PERCONA LIVE

Regression = git bisect



Regression = git bisect

- Manual:
 - git bisect start mysql-8.0.20 mysql-8.0.19
 - test
 - git bisect [good | bad]

Regression = git bisect

- Automated:

```
git bisect run sh -c '  
compile_mysql.sh  
if [ "$?" -ne "0" ]; then  
    exit 125  
fi  
.mysql-test/mtr bisect  
if [ "$?" -eq "0" ]; then  
    exit 0  
else  
    exit 1  
fi'
```



PERCONA LIVE

Regression = git bisect

- Bug#30460528: RENAME FIELD::REAL_MAYBE_NULL() TO FIELD::IS_NULLABLE()
- <https://github.com/mysql/mysql-server/commit/3039fac3969f7c1521863bfe1513631986d2b6bd>

Regression = git bisect



sql/field_conv.cc

Viewed

```
@@ -600,7 +600,7 @@ Copy_field::Copy_field(MEM_ROOT *mem_root, Item_field *item) :  
Copy_field() {  
  
600     m_from_field = item->field->  
       new_field(mem_root, item->field->  
       table);  
601     if (m_from_field == nullptr) return;  
602  
603 -     if (m_from_field->is_nullable()) {  
603 +     if (m_from_field->is_nullable() ||  
       m_from_field->table->is_nullable()) {  
604         // We need to allocate one extra  
         byte for null handling.  
605         uchar *ptr = mem_root->  
           ArrayAlloc<uchar>(m_from_field->  
           pack_length() + 1);  
606         m_to_field =
```

GDB / Coredump / PMP

Internal thread Deadlock bug



GDB

- GNU Debugger
- Works in various languages (C / C++ / Go / others)
- Can be used:
 - a. Remote
 - b. Live process - `gdb -p PID`
 - c. Starting a process - `gdb --args mysqld --datadir=.....`
 - d. Offline (coredump) - `gdb bin/mysqld core.xxxxx`

GDB

- Break point - stop the execution when a function is called.
- Condition break point - same as above but with condition (var1 == value)
- Watchpoints - same as break point but stop the execution when a variable is read / written / or both
- Next - execute the code until the next line.
- Continue - execute the code until the next break point.
- Step - enter the function.
- Bt - Mostrar Backtrace / Stack trace.
- Frame - pular para frame específico dentro da backtrace.

Coredump

- Snapshot / Dump of process memory
- Used alongside with GDB + binary of process
- Allows to check variables when the snapshot was collected
- Normally collected when process crash
- Can be collected on demand by gcore or gdb generate-core-file (eg: mysqld is frozen and I can't get in)

Coredump

- On crash:
 - Requires mysqld to be configured with `--core-file` and linux config:

```
echo 2 > /proc/sys/fs/suid_dumpable
mkdir /tmp/corefiles
chmod 777 /tmp/corefiles
echo "/tmp/corefiles/core" >
/proc/sys/kernel/core_pattern
echo "1" > /proc/sys/kernel/core_uses_pid
Mysqld --core-file
```
 - PS - mysqld --coredumper=/PATH

Coredump

```
-----  
5024 Thread 0x7f37ed6ef700 (LWP 82840) __lll_lock_wait (futex=futex@entry=0x55d6b662d1d8, private=0) at lowlevellock.c:52  
5025 Thread 0x7f37ed6a7700 (LWP 82841) __lll_lock_wait (futex=futex@entry=0x55d6b662d1d8, private=0) at lowlevellock.c:52  
5026 Thread 0x7f37ed65f700 (LWP 82846) __lll_lock_wait (futex=futex@entry=0x55d6b662d1d8, private=0) at lowlevellock.c:52  
5027 Thread 0x7f37ed617700 (LWP 82847) __lll_lock_wait (futex=futex@entry=0x55d6b662d1d8, private=0) at lowlevellock.c:52  
5028 Thread 0x7f37ed5cf700 (LWP 82848) __lll_lock_wait (futex=futex@entry=0x55d6b662d1d8, private=0) at lowlevellock.c:52  
5029 Thread 0x7f37ed587700 (LWP 82849) __lll_lock_wait (futex=futex@entry=0x55d6b662d1d8, private=0) at lowlevellock.c:52  
5030 Thread 0x7f37ed53f700 (LWP 82850) __lll_lock_wait (futex=futex@entry=0x55d6b662d1d8, private=0) at lowlevellock.c:52  
5031 Thread 0x7f37ed4f7700 (LWP 82851) __lll_lock_wait (futex=futex@entry=0x55d6b662d1d8, private=0) at lowlevellock.c:52  
5032 Thread 0x7f37ed4af700 (LWP 82852) __lll_lock_wait (futex=futex@entry=0x55d6b662d1d8, private=0) at lowlevellock.c:52  
5033 Thread 0x7f37ed467700 (LWP 82853) __lll_lock_wait (futex=futex@entry=0x55d6b662d1d8, private=0) at lowlevellock.c:52  
5034 Thread 0x7f37ed41f700 (LWP 82854) __lll_lock_wait (futex=futex@entry=0x55d6b662d1d8, private=0) at lowlevellock.c:52  
5035 Thread 0x7f37ed3d7700 (LWP 82855) __lll_lock_wait (futex=futex@entry=0x55d6b662d1d8, private=0) at lowlevellock.c:52  
5036 Thread 0x7f37ed38f700 (LWP 82856) __lll_lock_wait (futex=futex@entry=0x55d6b662d1d8, private=0) at lowlevellock.c:52  
5037 Thread 0x7f37ed347700 (LWP 82857) __lll_lock_wait (futex=futex@entry=0x55d6b662d1d8, private=0) at lowlevellock.c:52  
5038 Thread 0x7f37ed2ff700 (LWP 82858) __lll_lock_wait (futex=futex@entry=0x55d6b662d1d8, private=0) at lowlevellock.c:52
```

(gdb) █

Poor Man's Profiler - PMP

- Aggregates identical stack traces
- Very useful for core dumps with high number of threads
- <https://poormansprofiler.org/>
- Available via percona toolkit

pt-pmp

4996 __lll_lock_wait, __GI__pthread_mutex_lock, native_mutex_lock(thr_mutex.h:93), safe_mutex_lock(thr_mutex.cc:117), my_mutex_lock(thr_mutex.h:180), inline_mysql_mutex_lock(mysql_mutex.h:250), plugin_thdvar_init(sql_plugin.cc:2832), THD::init(sql_class.cc:740)

10 syscall, libaio::??(libaio.so.1), LinuxAIOHandler::collect(os0file.cc:2326), LinuxAIOHandler::poll(os0file.cc:2457), os_aio_linux_handler(os0file.cc:2506), os_aio_handler(os0file.cc:5820), fil_aio_wait(fil0fil.cc:7394), io_handler_thread(srv0start.cc:267)

3 futex_wait_cancelable, __pthread_cond_wait_common, __pthread_cond_wait, os_event::wait(os0event.cc:179), os_event::wait_low(os0event.cc:337), os_event_wait_low(os0event.cc:513), srv_worker_thread(srv0srv.cc:2668), void(srv0srv.cc:2668)

2 futex_abstimed_wait_cancelable, __pthread_cond_wait_common, __pthread_cond_timedwait, native_cond_timedwait(thr_cond.h:100), safe_cond_timedwait(thr_cond.cc:113), my_cond_timedwait(thr_cond.h:147), inline_mysql_cond_timedwait(mysql_cond.h:217), xpl::Cond::timed_wait(cond.cc:44)

1 syscall, epoll_wait(epoll_sub.c:64), epoll_dispatch(epoll.c:462), event_base_loop(event.c:1947), ngs::Socket_events::loop(socket_events.cc:173), ngs::Socket_acceptors_task::loop(socket_acceptors_task.cc:206), ngs::Server::run_task(server.cc:99), ngs::Server::start(server.cc:133)

1 futex_wait_cancelable, __pthread_cond_wait_common, __pthread_cond_wait, os_event::wait(os0event.cc:179), os_event::wait_low(os0event.cc:337), os_event_wait_low(os0event.cc:513), srv_purge_coordinator_suspend(srv0srv.cc:2798), srv_purge_coordinator_thread(srv0srv.cc:2912)

GDB / Coredump / pt-pmp

```
(gdb) t 5038
[Switching to thread 5038 (Thread 0x7fae2c147700 (LWP 899658))]
#0  __lll_lock_wait (futex=futex@entry=0x55f9c9d363a8, private=0) at lowlevellock.c:52
52      in lowlevellock.c
(gdb) bt 9
#0  __lll_lock_wait (futex=futex@entry=0x55f9c9d363a8, private=0) at lowlevellock.c:52
#1  0x00007fb2f1840235 in __GI___pthread_mutex_lock (mutex=0x55f9c9d363a8) at ../nptl/pthread_
mutex_lock.c:135
#2  0x000055f9c6ebb60a in native_mutex_lock (mutex=0x55f9c9d363a8) at /work/mysql/src/include/
thr_mutex.h:93
#3  0x000055f9c6ebb875 in safe_mutex_lock (mp=0x55f9c9d36380, try_lock=false, file=0x55f9c7f38
d48 "/work/mysql/src/sql/sql_plugin.cc", line=2832) at /work/mysql/src/mysys/thr_mutex.cc:117
#4  0x000055f9c59d9e1b in my_mutex_lock (mp=0x55f9c9606140 <LOCK_global_system_variables>, fil
e=0x55f9c7f38d48 "/work/mysql/src/sql/sql_plugin.cc", line=2832) at /work/mysql/src/include/th
r_mutex.h:180
#5  0x000055f9c59d9fd8 in inline_mysql_mutex_lock (that=0x55f9c9606140 <LOCK_global_system_var
iables>, src_file=0x55f9c7f38d48 "/work/mysql/src/sql/sql_plugin.cc", src_line=2832) at /work/
mysql/src/include/mysql/psi/mysql_mutex.h:250
#6  0x000055f9c59e3271 in plugin_thdvar_init (thd=0x7fb088058630, enable_plugins=true) at /wor
k/mysql/src/sql/sql_plugin.cc:2832
#7  0x000055f9c591ddd9 in THD::init (this=0x7fb088058630) at /work/mysql/src/sql/sql_class.cc:
740
#8  0x000055f9c591cbc6 in THD::THD (this=0x7fb088058630, enable_plugins=true) at /work/mysql/s
rc/sql/sql_class.cc:507
(More stack frames follow...)
(gdb) █
```

GDB / Coredump / pt-pmp

```
(gdb) f 2
#2 0x00005f9c6ebb60a in native_mutex_lock (mutex=0x55f9c9d363a8) at /work/mysql/src/include/
thr_mutex.h:93
93      return pthread_mutex_lock(mutex);
(gdb) p mutex->__data->__owner
$1 = 889019
(gdb) thread find 889019
Thread 38 has target id 'Thread 0x7fb2dc094700 (LWP 889019)'
```

GDB / Coredump / pt-pmp

- All 4996 threads wait on mutex from Thread 38

GDB / Coredump / pt-pmp

```
(gdb) t 38
[Switching to thread 38 (Thread 0x7fb2dc094700 (LWP 889019))]
#0  __lll_lock_wait (futex=futex@entry=0x55f9c9d36cc8, private=0) at lowlevellock.c:52
52      lowlevellock.c: No such file or directory.
(gdb) f 26
#26 0x000055f9c59a93d3 in do_command (thd=0x7fb22c00c880) at /work/mysql/src/sql/sql_parse.cc:1260
1260      return_value = dispatch_command(thd, &com_data, command);
(gdb) p thd->m_query_string->str
$5 = 0x7fb22c0e9f88 "SELECT * FROM performance_schema.session_variables WHERE VARIABLE_NAME LI
KE 'binlog_transaction_dependency_tracking'"
(gdb) f 2
#2 0x000055f9c6ebb60a in native_mutex_lock (mutex=0x55f9c9d36cc8) at /work/mysql/src/include/
thr_mutex.h:93
93      return pthread_mutex_lock(mutex);
(gdb) p mutex->__data->__owner
$6 = 889099
(gdb) thread find 889099
Thread 44 has target id 'Thread 0x7fb2bc7b2700 (LWP 889099)'
(gdb) ■
```

GDB / Coredump / pt-pmp

- All 4996 threads are trying to connect and waiting on mutex from Thread 38
- Thread 38
 - SELECT * FROM performance_schema.session_variables WHERE VARIABLE_NAME LIKE 'binlog_transaction_dependency_tracking'
 - Wait on mutex from Thread 44

GDB / Coredump / pt-pmp

```
(gdb) t 44
[Switching to thread 44 (Thread 0x7fb2bc7b2700 (LWP 889099))]
#0  __lll_lock_wait (futex=futex@entry=0x55f9c9d37a88, private=0) at lowlevellock.c:52
52      in lowlevellock.c
(gdb) f 11
#11 0x000055f9c59a93d3 in do_command (thd=0x7fb214000e80) at /work/mysql/src/sql/sql_parse.cc:
1260
1260      return_value = dispatch_command(thd, &com_data, command);
(gdb) p thd->m_query_string->str
$7 = 0x7fb214009bf8 "SHOW BINARY LOGS"
(gdb) f 2
#2 0x000055f9c6ebb60a in native_mutex_lock (mutex=0x55f9c9d37a88) at /work/mysql/src/include/
thr_mutex.h:93
93      return pthread_mutex_lock(mutex);
(gdb) p mutex->__data->__owner
$8 = 889097
(gdb) thread find 889097
Thread 42 has target id 'Thread 0x7fb2c4052700 (LWP 889097)'
(gdb) █
```

GDB / Coredump / pt-pmp

- All 4996 threads are trying to connect and waiting on mutex from Thread 38
- Thread 38
 - SELECT * FROM performance_schema.session_variables WHERE VARIABLE_NAME LIKE 'binlog_transaction_dependency_tracking'
 - Wait on mutex from Thread 44
- Thread 44
 - SHOW BINARY LOGS
 - Wait on mutex from Thread 42

GDB / Coredump / pt-pmp

```
(gdb) t 42
[Switching to thread 42 (Thread 0x7fb2c4052700 (LWP 889097))]
#0  __lll_lock_wait (futex=futex@entry=0x7fb22c0275b8, private=0) at lowlevellock.c:52
52      in lowlevellock.c
(gdb) f 18
#18 0x000055f9c59a93d3 in do_command (thd=0x7fb21c000e80) at /work/mysql/src/sql/sql_parse.cc:
1260
1260      return_value = dispatch_command(thd, &com_data, command);
(gdb) p thd->m_query_string->str
$9 = 0x7fb21c009bf8 "PURGE BINARY LOGS BEFORE NOW()"
(gdb) f 2
#2 0x000055f9c6ebb60a in native_mutex_lock (mutex=0x7fb22c0275b8) at /work/mysql/src/include/
thr_mutex.h:93
93      return pthread_mutex_lock(mutex);
(gdb) p mutex->__data->__owner
$10 = 889019
(gdb) thread find 889019
Thread 38 has target id 'Thread 0x7fb2dc094700 (LWP 889019)'
```

GDB / Coredump / pt-pmp

- All 4996 threads are trying to connect and waiting on mutex from Thread 38
- Thread 38
 - SELECT * FROM performance_schema.session_variables WHERE VARIABLE_NAME LIKE 'binlog_transaction_dependency_tracking'
 - Wait on mutex from Thread 44
- Thread 44
 - SHOW BINARY LOGS
 - Wait on mutex from Thread 42
- Thread 42
 - PURGE BINARY LOGS BEFORE NOW()
 - Wait on mutex from Thread 38

GDB / Coredump / pt-pmp

- Loop in [wait-for graph](#)
- Bug: Deadlock during purge_logs_before_date - [PS-1049 #91941](#)

Record & Replay

Xtrabackup incremental backup bug



Record & Replay

- Created by Mozilla - <https://rr-project.org/>
- Open Source - <https://github.com/rr-debugger/rr>
- Record once, replay deterministically as many times you want
- **Let you replay the code backwards**
- Run on a GDB session
- Hard to reproduce bugs
- Complex bugs

Record & Replay

- [Replay the Execution of MySQL With RR \(Record and Replay\)](#)
- [PXB-2180](#)
- PXB crash randomly after applying N incremental backups
- Not always the same stack trace - Crash on different parts of InnoDB
- Always trying to apply a record on same Space=>Page

```
#12 0x00000000015ad05f in recv_parse_or_apply_log_rec_body (type=MLOG_COMP_REC_INSERT,
ptr=0x7f2849150556 "\003K4G", '\377' <repeats 13 times>, end_ptr=0x7f2849150573 "",
space_id=<optimized out>, page_no=<optimized out>, block=0x7f2847d7da00, mtr=0x7f286857b4f0,
parsed_bytes=18446744073709551615) at /home/marcelo.altmann/percona-xtrabackup/storage/innobase/
log/log0recv.cc:2002
2002          ptr = page_cur_parse_insert_rec(FALSE, ptr, end_ptr, block, index, mtr);
(gdb) p block->page->id
+p block->page->id
$3 = {
  m_space = 4294967294,
  m_page_no = 5
}
```

Record & Replay

- Page layout diverted somehow between PXB and Server
- Crash on N th inc is a consequence of something wrong on N-M th inc.
- Run MySQL under RR
- Make a copy of all backups so we can re-run –prepare when the issue happens
- Read the LSNs for this same page before/after each backup prepare (od).
- Identify all changes to `m_space = 4294967294` & `m_page_no = 5` at mysqld.
- Got a reproducible backup - 6th incremental was crashing

Record & Replay

- `m_space = 4294967294` correspond to the MySQL data dictionary (`mysql.ibd`) – [dict0dict.h:1146](#)
- On disk page, LSN is stored at the 16th byte of the page and has a size of 8 bytes – [fil0types.h:66](#)
- Pages are written sequentially to disk, as an example, for the default 16k page size, from bytes 1 to 16384 will have the data for page 0, from byte 16385 to 32768 data from page 1, and so on.
- Frame is raw data of a page – [buf0buf.h:1358](#)

Record & Replay

```
$ od -j $((16384 * 5 + 16)) -N 8 -t x1 full/mysql.ibd
0240020 00 00 00 00 01 10 21 85
0240030
```

BF: 0x1102185

```
Breakpoint 1, buf_flush_note_modification (block=0x7fd2df4ad750,
69      ut_ad(!srv_read_only_mode ||

++rr-set-suppress-run-hook 1
(rr) p/x block->frame[16]@8
+p/x block->frame[16]@8
$1 =  {[0x0] = 0x0,
[0x1] = 0x0,
[0x2] = 0x0,
[0x3] = 0x0,
[0x4] = 0x1,
[0x5] = 0x10,
[0x6] = 0x21,
[0x7] = 0x85}

(rr)
```

Record & Replay

```
xtrabackup --prepare --apply-log-only --target-dir=full/...
...
Shutdown completed; log sequence number 17897577
Number of pools: 1
210402 17:46:29 completed OK!
```

BF: 0x1102185

```
$ od -j $((16384 * 5 + 16)) -N 8 -t x1 full/mysql.ibd
0240020 00 00 00 00 01 11 07 06
0240030
```

AF: 0x1100706

```
(rr) c
Breakpoint 1, buf_flush_note_modification (block=0x7fd2df4ad750,
69      ut_ad(!srv_read_only_mode ||

++rr-set-suppress-run-hook 1
(rr) p/x block->frame[16]@8
+p/x block->frame[16]@8
$16 = {[0x0] = 0x0,
[0x1] = 0x0,
[0x2] = 0x0,
[0x3] = 0x0,
[0x4] = 0x1,
[0x5] = 0x11,
[0x6] = 0x7,
[0x7] = 0x6}
(rr)
```

Record & Replay

```
xtrabackup --prepare --apply-log-only --target-dir=full/ --incremental-dir=inc1/
.
.
.
Shutdown completed; log sequence number 19082430
.
.
.
210402 18:12:20 completed OK!
.
$ od -j $((16384 * 5 + 16)) -N 8 -t x1 full/mysql.ibd
0240020 00 00 00 00 01 23 19 06
0240030
.
(rr) c
Breakpoint 1, buf_flush_note_modification (block=0x7fd2df4ad750, start_lsn=20262)
69      ut_ad(!srv_read_only_mode || 
++rr-set-suppress-run-hook 1
(rr) p/x block->frame[16]@8
+p/x block->frame[16]@8
$17 = {[0x0] = 0x0,
[0x1] = 0x0,
[0x2] = 0x0,
[0x3] = 0x0,
[0x4] = 0x1,
[0x5] = 0x23,
[0x6] = 0x19,
[0x7] = 0x6}
(rr)
```

BF: 0x1102185
AF: 0x1100706
I1: 0x1231906

Record & Replay

```
xtrabackup --prepare --apply-log-only --target-dir=full/ --incremental-dir=inc3/...
.
.
.
Shutdown completed; log sequence number 21455916
.
.
.
210402 18:18:25 completed OK!
.

$ od -j $((16384 * 5 + 16)) -N 8 -t x1 full/mysql.ibd
0240020 00 00 00 00 01 47 4d 3f
0240030

.

Breakpoint 1, buf_flush_note_modification (block=0x7fd2df4ad750, start_lsn=255294
69      ut_ad(!srv_read_only_mode ||

++rr-set-suppress-run-hook 1
(rr) p/x block->frame[16]@8
+p/x block->frame[16]@8
$19 = {[0x0] = 0x0,
[0x1] = 0x0,
[0x2] = 0x0,
[0x3] = 0x0,
[0x4] = 0x1,
[0x5] = 0x47,
[0x6] = 0x4d,
[0x7] = 0x3f}
(rr)
```

BF: 0x1102185

AF: 0x1100706

I1: 0x1231906

. . .

I3: 0x1474d3f

Record & Replay

```
xtrabackup --prepare --apply-log-only --target-dir=full/ --incremental-dir=inc4/  
.  
Shutdown completed; log sequence number 23044902  
.  
210402 18:24:00 completed OK!  
.  
$ od -j $((16384 * 5 + 16)) -N 8 -t x1 full/mysql.ibd  
0240020 00 00 00 00 01 5f a3 26  
0240030  
.  
.  
Breakpoint 1, buf_flush_note_modification (block=0x7fd2df4ad750, start_lsn=27218.  
69      ut_ad(!srv_read_only_mode ||  
++rr->suppress-run-hook 1  
(rr) p/x block->frame[16]@8  
+p/x block->frame[16]@8  
$242 =  {[0x0] = 0x0,  
[0x1] = 0x0,  
[0x2] = 0x0,  
[0x3] = 0x0,  
[0x4] = 0x1,  
[0x5] = 0x9f,  
[0x6] = 0x3f,  
[0x7] = 0xc9}  
(rr)
```

BF: 0x1102185

AF: 0x1100706

I1: 0x1231906

.. . .

I3: 0x1474d3f

B: 0x15fa326

S: 0x19f3fc9

Record & Replay

```
(rr) p block->frame
+p block->frame
$243 = (unsigned char *) 0x7fd2e0758000 "\327\064X["
(rr) watch *(unsigned char *) 0x7fd2e0758000
(rr) reverse-cont ←
Hardware watchpoint 2: *(unsigned char *) 0x7fd2e0758000
Old value = 80 'P'
New value = 43 '+'
0x0000000004c13903 in mach_write_to_4 (b=0x7fd2e0758000 "+k*\304", n=1353655015)
135  b[0] = static_cast<byte>(n >> 24);
++rr-set-suppress-run-hook 1
++rr-set-suppress-run-hook 1
(rr) p/x buf_flush_init_for_writing::block->frame[16]@8
+p/x buf_flush_init_for_writing::block->frame[16]@8
$12 =  {[0x0] = 0x0,
 [0x1] = 0x0,
 [0x2] = 0x0,
 [0x3] = 0x0,
 [0x4] = 0x1,
 [0x5] = 0x47,
 [0x6] = 0x4d,
 [0x7] = 0x3f}
(rr)
```

BF: 0x1102185

AF: 0x1100706

I1: 0x1231906

• • •

I3: 0x1474d3f

B: 0x15fa326

S: 0x19f3fc9



Are you **passionate**
about **Open Source**!?

We're looking for you!
Join us!



#RemoteWork

APPLY NOW: percona.com/careers